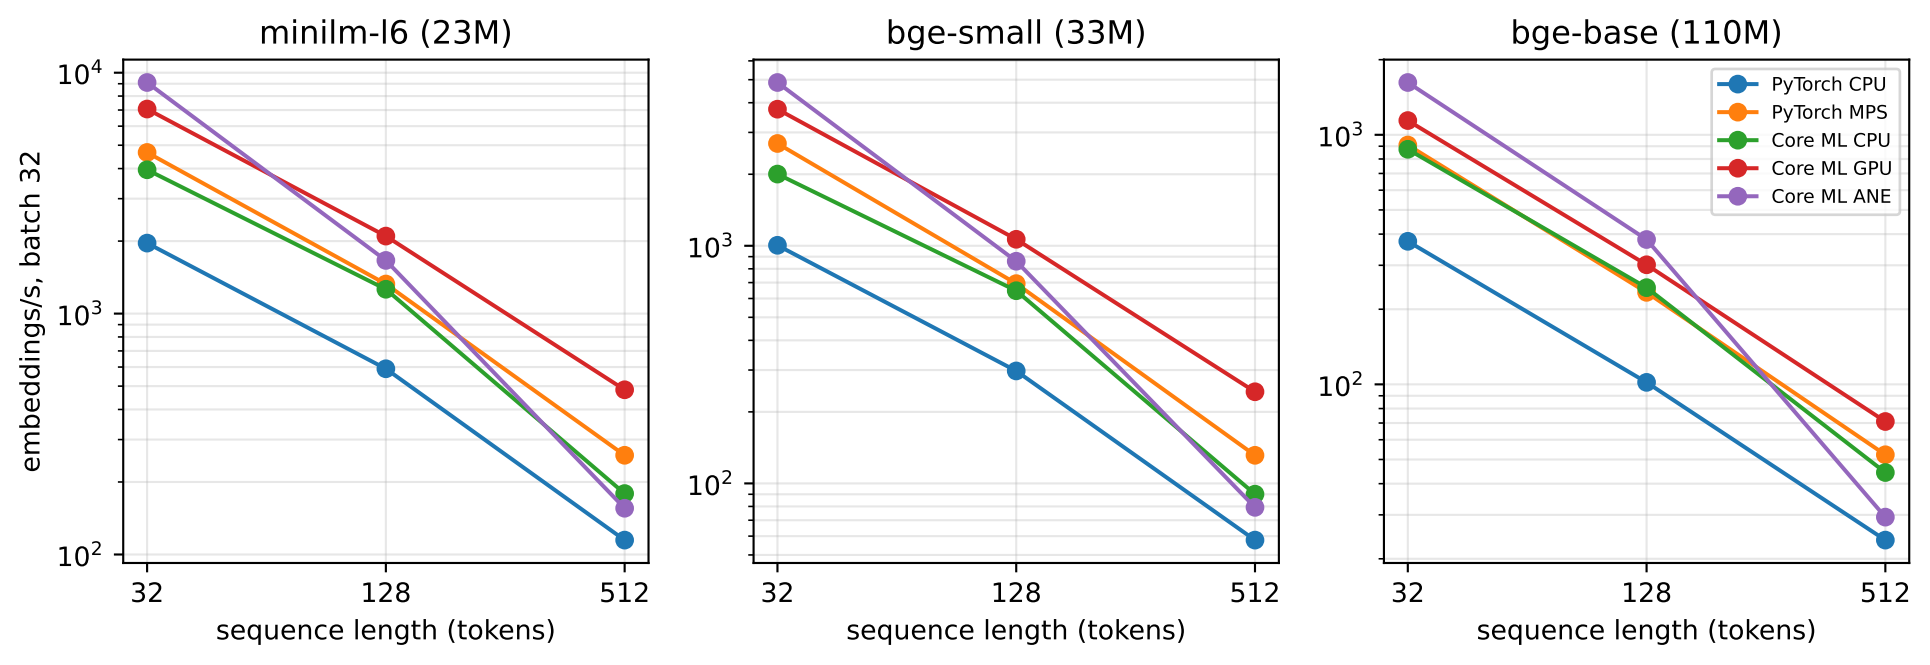
minilm-l6 (23M)
bge-small (33M)
bge-base (110M)
embeddings/s, batch 32
10^4
10^3
10^2
32
128
512
sequence length (tokens)
PyTorch CPU
PyTorch MPS
Core ML CPU
Core ML GPU
Core ML ANE